

ULTIMATE MySQL HANDBOOK



YOUR COMPLETE GUIDE TO
MySQL FROM BASICS TO
ADVANCED DATABASE MANAGEMENT

```
-- Create a database
CREATE DATABASE codewithharry;
USE codewithharry;
-- Create a table
CREATE TABLE users (
  id INT PRIMARY KEY AUTO_INCREMENT,
  name VARCHAR(100) NOT NULL,
  email VARCHAR(100) UNIQUE,
  created_at TIMESTAMP DEFAULT
    CURRENT_TIMESTAMP
);
-- Insert data
INSERT INTO users (name, email)
VALUES ('Harry', 'harry@cwah.dev');
-- Retrieve data
SELECT id, name, email
FROM users
WHERE id > 0;
```



CLEAR
EXPLANATIONS



PRACTICAL
EXAMPLES



REAL-WORLD
PROJECTS



BEST PRACTICES
AND TIPS



PERFECT FOR
BEGINNERS TO
ADVANCED



EASY TO
UNDERSTAND
CONCEPTS



HANDS-ON
SQL QUERIES



RELATIONAL
DATABASE
FUNDAMENTALS



OPTIMIZE AND
SCALE YOUR
DATABASES



SECURE YOUR
DATA AND
PERFORMANCE

CODE. PRACTICE. BUILD. REPEAT.

Installing MySQL

Beginner Friendly Database Guide

CHAPTER 01

What is MySQL Workbench?

MySQL Workbench is a visual tool used for database development and administration. It helps developers and database administrators design databases, write SQL queries, manage servers, create backups, and monitor database performance.

It provides an easy graphical interface for working with MySQL databases without relying completely on command line tools.

What is a Database Management System (DBMS)?

A Database Management System (DBMS) is software that allows users to create, store, organize, retrieve, and manage data efficiently.

It acts as a bridge between users and databases, making it easier to handle large amounts of information securely and systematically.

Features of MySQL

01

Open source and free to use database system.

02

High performance and reliable database engine.

03

Works on Windows, Linux, and macOS systems.

04

Provides strong security and user authentication features.

Windows / macOS Installation

1. Download from: <https://dev.mysql.com/downloads/installer/>
2. Run the installer and choose **Developer Default**
3. Follow the setup wizard instructions
4. Set a root password when prompted
5. Install MySQL Workbench for GUI-based management
6. Complete the installation and launch MySQL Workbench

Linux (Ubuntu) Installation

CHAPTER 02

MySQL Server Setup Guide

Installing MySQL on Ubuntu

Follow these steps to install MySQL Server and create a new database user.

Step 1: Update Package Index

TERMINAL COMMAND

```
sudo apt update
```

Step 2: Install MySQL Server

TERMINAL COMMAND

```
sudo apt install mysql-server
```

Step 3: Secure the Installation

TERMINAL COMMAND

```
sudo mysql_secure_installation
```

Choose your preferred security options during setup.

Step 4: Create a User 'harry'@'localhost'

Log into MySQL:

MYSQL LOGIN

```
sudo mysql
```

SQL User Commands

CHAPTER 02

Creating a MySQL User

Run the following SQL Commands

SQL COMMANDS

```
CREATE USER 'harry'@'localhost'  
IDENTIFIED BY 'password';  
  
GRANT ALL PRIVILEGES ON *.* TO  
'harry'@'localhost'  
WITH GRANT OPTION;  
  
FLUSH PRIVILEGES;  
  
EXIT;
```

Step 5: Test Login

TERMINAL COMMAND

```
mysql -u harry -p
```

Enter the password: password

Make sure to replace 'password' with a secure password of your choice in production environments.

Getting Started with MySQL

CHAPTER 03

Database Basics and Table Creation

What is a Database?

A database is a container that stores related data in an organized way. In MySQL, a database holds one or more tables.

Think of it like:

- **Folder analogy:**
 - A database is like a folder.
 - Each table is a file inside that folder.
 - The rows in the table are like the content inside each file.
- **Excel analogy:**
 - A database is like an Excel workbook.
 - Each table is a separate sheet inside that workbook.
 - Each row in the table is like a row in Excel.

Step 1: Create a Database

SQL COMMAND

```
CREATE DATABASE startersql;
```

After creating the database, either:

- Right-click it in MySQL Workbench and select **"Set as Default Schema"**
- Use this SQL command:

SQL COMMAND

```
USE startersql;
```

Creating Tables in MySQL

CHAPTER 03

Understanding Tables and Data Types

Step 2: Create a Table

Now we'll create a simple users table:

SQL COMMAND

```
CREATE TABLE users (  
  
  id INT AUTO_INCREMENT PRIMARY KEY,  
  
  name VARCHAR(100) NOT NULL,  
  
  email VARCHAR(100) UNIQUE NOT NULL,  
  
  gender ENUM('Male', 'Female', 'Other'),  
  
  date_of_birth DATE,  
  
  created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
  
);
```

Step 3: Drop the Database

You can delete the entire database (and all its tables) using:

SQL COMMAND

```
DROP DATABASE startersql;
```

Data Types Explained

- INT : Integer type, used for whole numbers.
- VARCHAR(100) : Variable-length string, up to 100 characters.
- ENUM : A string object with a value chosen from a list of permitted values.
- DATE : Stores date values.

Data Types & Constraints

CHAPTER 03

Understanding MySQL Data Structures

More Data Types

- **TIMESTAMP** : Stores date and time, automatically set to the current timestamp.
- **BOOLEAN** : Stores TRUE or FALSE values.
- **DECIMAL(10, 2)** : Stores exact numeric data values, useful for financial data. The first number is the total number of digits, and the second is the number of digits after the decimal point.

Constraints Explained

- **AUTO_INCREMENT** : Automatically generates a unique number for each row.
- **PRIMARY KEY** : Uniquely identifies each row in the table.
- **NOT NULL** : Ensures a column cannot have a NULL value.
- **UNIQUE** : Ensures all values in a column are different.
- **DEFAULT** : Sets a default value for a column if no value is provided.

EXAMPLE

```
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
  
is_active BOOLEAN DEFAULT TRUE
```

Working with Tables in MySQL

CHAPTER 04

Selecting and Renaming Tables

Selecting Data from a Table

Select All Columns

SQL COMMAND

```
SELECT * FROM users;
```

This fetches every column and every row from the users table.

Select Specific Columns

SQL COMMAND

```
SELECT name, email FROM users;
```

This only fetches the name and email columns from all rows.

Renaming a Table

To rename an existing table:

SQL COMMAND

```
RENAME TABLE users TO customers;
```

To rename it back:

SQL COMMAND

```
RENAME TABLE customers TO users;
```


Altering Tables in MySQL

CHAPTER 04

Modifying Existing Table Structure

Altering a Table

You can use ALTER TABLE to modify an existing table.

Add a Column

SQL COMMAND

```
ALTER TABLE users  
ADD COLUMN is_active BOOLEAN DEFAULT TRUE;
```

Drop a Column

SQL COMMAND

```
ALTER TABLE users  
DROP COLUMN is_active;
```

Modify a Column Type

SQL COMMAND

```
ALTER TABLE users  
MODIFY COLUMN name VARCHAR(150);
```

Move a Column to the First Position

To move a column (e.g., email) to the first position:

SQL COMMAND

```
ALTER TABLE users  
MODIFY COLUMN email VARCHAR(100) FIRST;
```

Advanced Table Modifications

CHAPTER 04

Managing Column Position and Structure

Move a Column After Another Column

To move a column after another column (e.g., move gender after name):

SQL COMMAND

```
ALTER TABLE users
MODIFY COLUMN gender ENUM('Male', 'Female', 'Other')
AFTER name;
```

Add Multiple Columns

You can also add multiple columns in a single query.

SQL COMMAND

```
ALTER TABLE users

ADD COLUMN phone VARCHAR(15),

ADD COLUMN city VARCHAR(100);
```

Rename a Column

Use RENAME COLUMN to change a column name.

SQL COMMAND

```
ALTER TABLE users
RENAME COLUMN phone TO mobile_number;
```

Inserting Data into MySQL Tables

CHAPTER 05

Adding Records Using INSERT INTO

To add data into a table, we use the `INSERT INTO` statement.

Insert Without Specifying Column Names (Full Row Insert)

This method requires you to provide values for all columns in order, except columns with default values or `AUTO_INCREMENT`.

SQL COMMAND

```
INSERT INTO users VALUES  
  
(1, 'Alice', 'alice@example.com',  
 'Female', '1995-05-14', DEFAULT);
```

Not recommended if your table structure might change later (e.g., new columns added later).

Insert by Specifying Column Names (Best Practice)

This method is safer and more readable. You only insert into specific columns.

SQL COMMAND

```
INSERT INTO users  
(name, email, gender, date_of_birth)  
  
VALUES  
  
('Bob', 'bob@example.com',  
 'Male', '1990-11-23');
```

Multiple Row Inserts

CHAPTER 05

Efficient Data Insertion Techniques

Insert Multiple Rows at Once

You can insert multiple rows in a single query.

SQL COMMAND

```
INSERT INTO users  
(name, email, gender, date_of_birth)
```

VALUES

```
( 'Bob', 'bob@example.com',  
  'Male', '1990-11-23'),  
  
( 'Charlie', 'charlie@example.com',  
  'Other', '1988-02-17');
```

The remaining columns like `id` (which is `AUTO_INCREMENT`) and `created_at` (which has a default value) are automatically handled by MySQL.

Efficient Bulk Inserts

CHAPTER 05

Adding Multiple Records in One Query

Writing multiple rows inside a single `INSERT INTO` statement improves performance and reduces the number of database queries.

This approach is commonly used while importing large datasets or inserting multiple users at once.

Insert Multiple Rows Efficiently

SQL COMMAND

```
INSERT INTO users  
(name, email, gender, date_of_birth)
```

VALUES

```
('Charlie', 'charlie@example.com',  
'Other', '1988-02-17'),  
  
( 'David', 'david@example.com',  
'Male', '2000-08-09'),  
  
( 'Eva', 'eva@example.com',  
'Female', '1993-12-30');
```


Querying Data in MySQL using SELECT

CHAPTER 06

Retrieving and Filtering Data

The SELECT statement is used to query data from a table.

Basic Syntax

SQL SYNTAX

```
SELECT column1, column2 FROM table_name;
```

To select all columns:

SQL COMMAND

```
SELECT * FROM users;
```

Filtering Rows with WHERE

Equal To

SQL COMMAND

```
SELECT * FROM users  
WHERE gender = 'Male';
```



Advanced WHERE Conditions

CHAPTER 06

Working with Comparisons and NULL

Not Equal To

SQL COMMAND

```
SELECT * FROM users
WHERE gender != 'Female';
```

-- OR

```
SELECT * FROM users
WHERE gender <> 'Female';
```

Greater Than / Less Than

SQL COMMAND

```
SELECT * FROM users
WHERE date_of_birth < '1995-01-01';
```

```
SELECT * FROM users
WHERE id > 10;
```

Greater Than or Equal / Less Than or Equal

SQL COMMAND

```
SELECT * FROM users
WHERE id >= 5;
```

```
SELECT * FROM users
WHERE id <= 20;
```

NULL & BETWEEN

CHAPTER 06

Handling Missing Values and Ranges

Working with NULL

IS NULL

SQL COMMAND

```
SELECT * FROM users
WHERE date_of_birth IS NULL;
```

IS NOT NULL

SQL COMMAND

```
SELECT * FROM users
WHERE date_of_birth IS NOT NULL;
```

BETWEEN

SQL COMMAND

```
SELECT * FROM users
WHERE date_of_birth BETWEEN
'1990-01-01' AND '2000-12-31';
```



Pattern Matching & Sorting

CHAPTER 06

LIKE, AND / OR and ORDER BY

IN

SQL COMMAND

```
SELECT * FROM users
WHERE gender IN ('Male', 'Other');
```

LIKE (Pattern Matching)

SQL COMMAND

```
SELECT * FROM users
WHERE name LIKE 'A%';

-- Starts with A

SELECT * FROM users
WHERE name LIKE '%a';

-- Ends with a

SELECT * FROM users
WHERE name LIKE '%li%';

-- Contains 'li'
```

AND / OR

SQL COMMAND

```
SELECT * FROM users
WHERE gender = 'Female'
AND date_of_birth > '1990-01-01';

SELECT * FROM users
WHERE gender = 'Male'
OR gender = 'Other';
```

ORDER BY & LIMIT

CHAPTER 06

Sorting and Limiting Query Results

ORDER BY

SQL COMMAND

```
SELECT * FROM users
ORDER BY date_of_birth ASC;
```

```
SELECT * FROM users
ORDER BY name DESC;
```

LIMIT

SQL COMMAND

```
SELECT * FROM users
LIMIT 5;
```

-- Top 5 rows

```
SELECT * FROM users
LIMIT 10 OFFSET 5;
```

-- Skip first 5 rows

```
SELECT * FROM users
LIMIT 5, 10;
```

-- Get 10 rows

-- starting from the 6th row

SQL COMMAND

```
SELECT * FROM users
ORDER BY created_at DESC
LIMIT 10;
```



Quick Quiz

CHAPTER 06

Practice SQL Query Reading

What does the following queries do?

QUERY 1

```
SELECT * FROM users
WHERE salary > 60000
ORDER BY created_at DESC
LIMIT 5;
```

QUERY 2

```
SELECT * FROM users
ORDER BY salary DESC;
```

QUERY 3

```
SELECT * FROM users
WHERE salary BETWEEN
50000 AND 70000;
```



UPDATE - Modifying Existing Data

CHAPTER 07

Changing Values in Existing Rows

The UPDATE statement is used to change values in one or more rows.

Basic Syntax

SQL SYNTAX

```
UPDATE table_name
SET column1 = value1, column2 = value2
WHERE condition;
```

Example: Update One Column

SQL COMMAND

```
UPDATE users
SET name = 'Alicia'
WHERE id = 1;
```

This changes the name of the user with `id = 1` to "Alicia".

Example: Update Multiple Columns

SQL COMMAND

```
UPDATE users
SET name = 'Robert', email = 'robert@example.com'
WHERE id = 2;
```

Without WHERE Clause (Warning)

SQL COMMAND

```
UPDATE users
SET gender = 'Other';
```

This updates every row in the table. Be very careful when omitting the WHERE clause.

Quick Quiz: Practice Your UPDATE Skills

CHAPTER 07

Writing Real SQL Update Queries

Try answering or running these queries based on your users table.

1. Update the salary of user with id = 5 to ₹70,000.

SQL COMMAND

```
UPDATE users
SET salary = 70000
WHERE id = 5;
```

2. Change the name of the user with email aisha@example.com to Aisha Khan.

SQL COMMAND

```
UPDATE users
SET name = 'Aisha Khan'
WHERE email = 'aisha@example.com';
```

3. Increase salary by ₹10,000 for all users whose salary is less than ₹60,000.

SQL COMMAND

```
UPDATE users
SET salary = salary + 10000
WHERE salary < 60000;
```

4. Set the gender of user Ishaan to Other.

SQL COMMAND

```
UPDATE users
SET gender = 'Other'
WHERE name = 'Ishaan';
```



More UPDATE Practice Queries

CHAPTER 07

Understanding Row Modifications

5. Reset salary of all users to ₹50,000 (Careful - affects all rows).

SQL COMMAND

```
UPDATE users  
SET salary = 50000;
```

Note: This query will overwrite salary for every user. Use with caution!



DELETE - Removing Data from a Table

CHAPTER 08

Deleting Rows from Tables

The DELETE statement removes rows from a table.

Basic Syntax

SQL SYNTAX

```
DELETE FROM table_name  
WHERE condition;
```

Example: Delete One Row

SQL COMMAND

```
DELETE FROM users  
WHERE id = 3;
```

Delete Multiple Rows

SQL COMMAND

```
DELETE FROM users  
WHERE gender = 'Other';
```

Advanced DELETE Operations

CHAPTER 08

Removing Data Safely

Delete All Rows (but keep table structure)

SQL COMMAND

```
DELETE FROM users;
```

Drop the Entire Table (use with caution)

SQL COMMAND

```
DROP TABLE users;
```

This removes the table structure and all data permanently.

Best Practices

- Always use `WHERE` unless you're intentionally deleting everything.
- Run a `SELECT` query with the same condition first before deleting data.
- Always back up important data before performing destructive operations.

SQL CHECK QUERY

```
SELECT * FROM users  
WHERE id = 3;
```



Quick Quiz: Practice DELETE Queries

CHAPTER 08

Understand What These Queries Will Do

What will happen if you run the following queries?

QUERY 1

```
DELETE FROM users
WHERE salary < 50000;
```

QUERY 2

```
DELETE FROM users
WHERE salary IS NULL;
```

QUERY 3

```
DELETE FROM users
WHERE gender = 'Male'
AND salary > 100000;
```

QUERY 4

```
DELETE FROM users
WHERE created_at > '2025-01-01';
```



MySQL Constraints

CHAPTER 09

Constraints in MySQL are rules applied to table columns to ensure the accuracy, validity, and integrity of the data.

1. UNIQUE Constraint

Ensures that all values in a column are **different**.

Example (during table creation):

```
CREATE TABLE users (  
    id INT PRIMARY KEY,  
    email VARCHAR(100) UNIQUE  
);
```

Add UNIQUE using ALTER TABLE:

```
ALTER TABLE users  
ADD CONSTRAINT unique_email  
UNIQUE (email);
```


2. NOT NULL Constraint

Ensures that a column **cannot contain NULL** values.

Example:

```
CREATE TABLE users (  
  
    id INT PRIMARY KEY,  
  
    name VARCHAR(100) NOT NULL  
  
);
```

Change an existing column to NOT NULL:

```
ALTER TABLE users  
  
MODIFY COLUMN name  
VARCHAR(100) NOT NULL;
```

Make a column nullable again:

```
ALTER TABLE users  
  
MODIFY COLUMN name  
VARCHAR(100) NULL;
```

3. CHECK Constraint

Ensures that values in a column satisfy a **specific condition**.

Example: Allow only dates of birth after Jan 1, 2000

```
ALTER TABLE users

ADD CONSTRAINT chk_dob
CHECK (date_of_birth > '2000-01-01');
```

Naming the constraint (chk_dob) helps if you want to drop it later.

4. DEFAULT Constraint

Sets a default value for a column if none is provided during insert.

Example:

```
CREATE TABLE users (

    id INT PRIMARY KEY,

    is_active BOOLEAN DEFAULT TRUE

);
```

Add DEFAULT using ALTER TABLE:

```
ALTER TABLE users

ALTER COLUMN is_active
SET DEFAULT TRUE;
```

5. PRIMARY KEY Constraint

Uniquely identifies each row. Must be **NOT NULL** and **UNIQUE**.

Example:

```
CREATE TABLE users (  
    id INT PRIMARY KEY,  
    name VARCHAR(100)  
);
```

Add later with ALTER TABLE:

```
ALTER TABLE users  
ADD PRIMARY KEY (id);
```



6. AUTO_INCREMENT

Used with PRIMARY KEY to automatically assign the next number.

Example:

```
CREATE TABLE users (  
  
    id INT AUTO_INCREMENT PRIMARY KEY,  
  
    name VARCHAR(100)  
  
);
```

Summary Table

Constraint	Purpose
UNIQUE	Prevents duplicate values
NOT NULL	Ensures value is not NULL
CHECK	Restricts values using a condition
DEFAULT	Sets a default value
PRIMARY KEY	Uniquely identifies each row
AUTO_INCREMENT	Automatically generates unique numbers

SQL Functions (MySQL)

CHAPTER 10

Aggregate and String Functions

SQL functions help you analyze, transform, or summarize data in your tables.

We'll use the users table which includes:

id, name, email, gender, date_of_birth, salary, created_at

1. Aggregate Functions

COUNT()

Count total number of users:

```
SELECT COUNT(*) FROM users;
```

Count users who are Female:

```
SELECT COUNT(*) FROM users  
WHERE gender = 'Female';
```

MIN() and MAX()

Get the minimum and maximum salary:

```
SELECT MIN(salary) AS min_salary,  
MAX(salary) AS max_salary  
FROM users;
```

Aggregate & String Functions

CHAPTER 10

Data Analysis Functions

SUM()

Calculate total salary payout:

```
SELECT SUM(salary) AS total_payroll
FROM users;
```

AVG()

Find average salary:

```
SELECT AVG(salary) AS avg_salary
FROM users;
```

Grouping with GROUP BY

Average salary by gender:

```
SELECT gender,
AVG(salary) AS avg_salary
FROM users
GROUP BY gender;
```

String Functions

CHAPTER 10

Working with Text Data

2. String Functions

LENGTH()

Length of user names:

```
SELECT name,  
LENGTH(name) AS name_length  
FROM users;
```

LOWER() and UPPER()

Convert names to lowercase or uppercase:

```
SELECT name,  
LOWER(name) AS lowercase_name  
FROM users;  
  
SELECT name,  
UPPER(name) AS uppercase_name  
FROM users;
```

CONCAT()

Combine name and email:

```
SELECT CONCAT(name, ' <', email, '>')  
AS user_contact  
FROM users;
```


Date Functions

CHAPTER 10

Working with Dates and Time

3. Date Functions

NOW()

Current date and time:

```
SELECT NOW();
```

YEAR(), MONTH(), DAY()

Extract parts of date_of_birth:

```
SELECT name,  
YEAR(date_of_birth) AS birth_year  
FROM users;
```

DATEDIFF()

Find number of days between today and birthdate:

```
SELECT name,  
DATEDIFF(CURDATE(), date_of_birth)  
AS days_lived  
FROM users;
```

Date & Mathematical Functions

CHAPTER 10

Advanced Calculations in SQL

TIMESTAMPDIFF()

Calculate age in years:

```
SELECT name,  
TIMESTAMPDIFF(YEAR, date_of_birth, CURDATE())  
AS age  
FROM users;
```

4. Mathematical Functions

ROUND(), FLOOR(), CEIL()

```
SELECT salary,  
  
ROUND(salary) AS rounded,  
  
FLOOR(salary) AS floored,  
  
CEIL(salary) AS ceiled  
  
FROM users;
```

MOD()

Find even or odd user IDs:

```
SELECT id,  
MOD(id, 2) AS remainder  
FROM users;
```

Conditional Functions

CHAPTER 10

Logical Conditions in SQL

5. Conditional Functions

IF()

```
SELECT name,  
gender,  
  
IF(gender = 'Female', 'Yes', 'No')  
AS is_female  
  
FROM users;
```

Summary Table

Function	Purpose
COUNT()	Count rows
SUM()	Total of a column
AVG()	Average of values
MIN() / MAX()	Lowest / highest value
LENGTH()	String length
CONCAT()	Merge strings
YEAR() / DATEDIFF()	Date breakdown / age
ROUND()	Rounding numbers
IF()	Conditional logic

MySQL Transactions and AutoCommit

CHAPTER 11

Managing Transactions in MySQL

By default, MySQL operates in **AutoCommit** mode. This means that every SQL statement is treated as a **transaction** and is committed automatically. However, for more control over when changes are saved, you can turn AutoCommit off and manage transactions manually.

1. Disabling AutoCommit

When **AutoCommit** is off, you can explicitly control when to commit or rollback changes.

To disable AutoCommit:

```
SET autocommit = 0;
```

This turns off AutoCommit, meaning that changes you make won't be saved to the database unless you explicitly tell MySQL to commit them.

Important: Until you execute a `COMMIT`, your changes are not permanent.

2. COMMIT: Save Changes to the Database

Once you've made changes and you're confident that everything is correct, you can use the `COMMIT` command to save those changes.

To commit a transaction:

```
COMMIT;
```

This saves all the changes made since the last `COMMIT` or `ROLLBACK`. After this point, the changes become permanent.

Rollback and Transactions

CHAPTER 11

Reverting Database Changes Safely

3. ROLLBACK: Revert Changes to the Last Safe Point

If you make an error or decide you don't want to save your changes, you can **rollback** the transaction to its previous state.

To rollback a transaction:

```
ROLLBACK;
```

This undoes all changes since the last `COMMIT` or `ROLLBACK`.

Example Workflow

1. Turn off AutoCommit:

```
SET autocommit = 0;
```

2. Make some changes (e.g., updating a salary):

```
UPDATE users  
SET salary = 80000  
WHERE id = 5;
```

3. Decide whether to commit or rollback:

1. If you are happy with the changes, run:

```
COMMIT;
```

2. If you are not happy and want to revert the changes, run:

```
ROLLBACK;
```

AutoCommit Best Practices

CHAPTER 11

Working Safely with Transactions

4. Enabling AutoCommit Again

If you want to turn AutoCommit back on (so that every statement is automatically committed), you can do so with:

```
SET autocommit = 1;
```

Best Practices

- Use `COMMIT` when you want to make changes permanent.
- Use `ROLLBACK` to discard changes if something goes wrong.
- Consider **disabling AutoCommit** when performing complex updates to avoid saving partial or incorrect data.
- Always test transaction queries carefully before applying them in production databases.

Understanding PRIMARY KEY in MySQL

CHAPTER 12

Primary Keys and Unique Constraints

A **PRIMARY KEY** is a constraint in SQL that **uniquely identifies each row** in a table. It is one of the most important concepts in database design.

What is a Primary Key?

- A **PRIMARY KEY** :
 - Must be **unique**
 - Cannot be **NULL**
 - Is used to identify rows in a table
 - Can be a single column or a combination of columns
 - Each table can have only **one** primary key

Example:

```
CREATE TABLE users (  
    id INT AUTO_INCREMENT PRIMARY KEY,  
    name VARCHAR(100)  
);
```

UNIQUE and Constraints

CHAPTER 12

Comparing PRIMARY KEY and UNIQUE

How Is PRIMARY KEY Different from UNIQUE ?

At first glance, PRIMARY KEY and UNIQUE might seem similar since both prevent duplicate values. But there are important differences:

Feature	PRIMARY KEY	UNIQUE
Must be unique	Yes	Yes
Allows NULL values	No	Yes (one or more NULLs allowed)
How many allowed	Only one per table	Can have multiple
Required by table	Recommended, often required	Optional
Dropping	Cannot be easily dropped	Can be dropped anytime

Example with UNIQUE

```
CREATE TABLE users (  
    id INT AUTO_INCREMENT PRIMARY KEY,  
    email VARCHAR(100) UNIQUE,  
    name VARCHAR(100)  
);
```

In this example:

- `id` is the unique identifier for each row.
- `email` must be unique, but is not the primary key.

Dropping Keys

CHAPTER 12

Managing Constraints in MySQL

Can I Drop a PRIMARY KEY?

Yes, but it is more restricted than dropping a `UNIQUE` constraint.

```
ALTER TABLE users  
DROP PRIMARY KEY;
```

This may fail if the primary key is being used elsewhere (like in a foreign key or auto_increment column).

To drop a `UNIQUE` constraint:

```
ALTER TABLE users  
DROP INDEX email;
```



Auto Increment and Takeaways

CHAPTER 12

Best Practices with Primary Keys

Auto Increment

In MySQL, a `PRIMARY KEY` is often used with the `AUTO_INCREMENT` attribute to automatically generate unique values for new rows.

```
CREATE TABLE users (  
    id INT AUTO_INCREMENT PRIMARY KEY,  
    name VARCHAR(100)  
);
```

This means that every time you insert a new row, MySQL will automatically assign a unique value to the `id` column.

```
ALTER TABLE users  
AUTO_INCREMENT = 1000;
```

Key Takeaways

- Use `PRIMARY KEY` for the **main identifier** of a row.
- Use `UNIQUE` for enforcing **non-duplicate values** in other columns (like email or phone).
- You can have only **one** primary key, but you can have **many unique constraints**.
- Combine `AUTO_INCREMENT` with `PRIMARY KEY` for automatically generated IDs.
- Always choose a stable and unique column for your primary key.

Foreign Keys in MySQL

CHAPTER 13

Understanding Relationships Between Tables

A foreign key is a column that creates a **link between two tables**. It ensures that the value in one table **must match** a value in another table.

This is used to maintain **data integrity** between related data.

Why Use Foreign Keys?

Imagine this scenario:

You have a `users` table. Now you want to store each user's address. Instead of putting address columns inside the `users` table, you create a separate `addresses` table, and link it to `users` using a foreign key.

Creating a Table with a Foreign Key

Let's create an `addresses` table where each address belongs to a user.

```
CREATE TABLE addresses (  
    id INT AUTO_INCREMENT PRIMARY KEY,  
    user_id INT,  
    street VARCHAR(255),  
    city VARCHAR(100),  
    state VARCHAR(100),  
    pincode VARCHAR(10),  
    FOREIGN KEY (user_id)  
    REFERENCES users(id)  
);
```

Managing Foreign Keys

CHAPTER 13

Creating and Removing Constraints

Explanation

- `user_id` is a foreign key.
- It references the `id` column in the `users` table.
- This ensures that every address must be linked to a valid user.

Dropping a Foreign Key

To drop a foreign key, you need to know its **constraint name**. MySQL auto-generates it if you don't specify one, or you can name it yourself.

```
CREATE TABLE addresses (  
    id INT AUTO_INCREMENT PRIMARY KEY,  
    user_id INT,  
    CONSTRAINT fk_user  
    FOREIGN KEY (user_id)  
    REFERENCES users(id)  
);
```

To drop it:

```
ALTER TABLE addresses  
DROP FOREIGN KEY fk_user;
```



ALTER Operations

Adding Foreign Keys Later

Adding a Foreign Key Later (Using ALTER)

Suppose the foreign key was **not defined during table creation**. You can add it later using `ALTER TABLE` :

```
ALTER TABLE addresses
```

```
ADD CONSTRAINT fk_user
```

```
FOREIGN KEY (user_id)
```

```
REFERENCES users(id);
```



ON DELETE Actions

CHAPTER 13

Controlling Delete Behavior

Adding ON DELETE Action

By default, if you delete a user that has related addresses, MySQL will throw an error. You can control this behavior with `ON DELETE`.

Example with `ON DELETE CASCADE` :

If you want addresses to be automatically deleted when the user is deleted:

```
CREATE TABLE addresses (  
    id INT AUTO_INCREMENT PRIMARY KEY,  
    user_id INT,  
    street VARCHAR(255),  
    city VARCHAR(100),  
    state VARCHAR(100),  
    pincode VARCHAR(10),  
    CONSTRAINT fk_user  
    FOREIGN KEY (user_id)  
    REFERENCES users(id)  
    ON DELETE CASCADE  
);
```

Foreign Key Summary

CHAPTER 13

Delete Options and Final Notes

Or alter it later:

```
ALTER TABLE addresses  
  
ADD CONSTRAINT fk_user  
FOREIGN KEY (user_id)  
REFERENCES users(id)  
ON DELETE CASCADE;
```

Other ON DELETE Options

ON DELETE Option	Behavior
CASCADE	Deletes all related rows in child table
SET NULL	Sets the foreign key to NULL in the child table
RESTRICT	Prevents deletion of parent if child exists (default)

Summary

- Foreign keys **connect tables** and enforce **valid references**.
- You can create them inline or with `ALTER TABLE`.
- You can drop them by name.
- Use `ON DELETE` to control what happens when the parent row is deleted.

SQL JOINS in MySQL

CHAPTER 14

Combining Data from Multiple Tables

In SQL, JOINS are used to combine rows from **two or more tables** based on related columns. Usually a **foreign key** in one table references a **primary key** in another.

We'll use the following two tables:

users table

id	name
1	Aarav
2	Sneha
3	Raj

addresses table

id	user_id	city
1	1	Mumbai
2	2	Kolkata
3	4	Delhi

Note: `user_id` is a foreign key that references `users.id`.

INNER JOIN

Matching Rows from Both Tables

1. INNER JOIN

Returns **only the matching rows** from both tables.

```
SELECT users.name, addresses.city  
  
FROM users  
  
INNER JOIN addresses  
ON users.id = addresses.user_id;
```

Output:

name	city
Aarav	Mumbai
Sneha	Kolkata

Raj is excluded because there is no matching address. Delhi is excluded because its user_id (4) is not in users .

Visual Representation:

```
users      addresses  
-----  
  
| 1 |      | 1 |  
| 2 |      | 2 |  
| 3 |      |  |  
  
=> only matching pairs
```

LEFT JOIN

All Left Rows with Matching Right Rows

2. LEFT JOIN

Returns **all rows from the left table** (`users`), and matching rows from the right table (`addresses`). If no match is found, NULLs are returned.

```
SELECT users.name, addresses.city  
  
FROM users  
  
LEFT JOIN addresses  
ON users.id = addresses.user_id;
```

Output:

name	city
Aarav	Mumbai
Sneha	Kolkata
Raj	NULL

`Raj` is shown even though he doesn't have an address.

Visual Representation:

users	addresses
----	-----
1	1
2	2
3	

=> all users + matched addresses (or NULL)

RIGHT JOIN and Summary

CHAPTER 14

Matching Right Side Data

3. RIGHT JOIN

Returns **all rows from the right table** (`addresses`), and matching rows from the left table (`users`). If no match is found, NULLs are returned.

```
SELECT users.name, addresses.city  
  
FROM users  
  
RIGHT JOIN addresses  
ON users.id = addresses.user_id;
```

Output:

name	city
Aarav	Mumbai
Sneha	Kolkata
NULL	Delhi

Delhi is shown even though it points to a `user_id` that doesn't exist.

Visual Representation:

```
users      addresses  
-----  
  
| 1 |      | 1 |  
| 2 |      | 2 |  
|   |      | 4 |
```

=> all addresses + matched users (or NULL)

Summary Table

SQL UNION and UNION ALL

CHAPTER 15

Combining Results from Multiple Queries

The `UNION` operator in SQL is used to **combine the result sets of two or more `SELECT` statements**. It removes duplicates by default.

If you want to include all rows including duplicates, use `UNION ALL`.

Example Scenario

You already have a `users` table for active users. Now, we'll create an `admin_users` table to store users who are administrators or have special roles.

We will then combine the names from both tables using `UNION`.

Step 1: Create the `admin_users` Table

```
CREATE TABLE admin_users (  
    id INT PRIMARY KEY,  
    name VARCHAR(100),  
    email VARCHAR(100),  
    gender ENUM('Male', 'Female', 'Other'),  
    date_of_birth DATE,  
    salary INT  
);
```

Using UNION

CHAPTER 15

Combining Query Results

Step 2: Insert Sample Data into admin_users

```
INSERT INTO admin_users
(id, name, email, gender, date_of_birth, salary)
VALUES

(101, 'Anil Kumar', 'anil@example.com',
'Male', '1985-04-12', 60000),

(102, 'Pooja Sharma', 'pooja@example.com',
'Female', '1992-09-20', 58000),

(103, 'Rakesh Yadav', 'rakesh@example.com',
'Male', '1989-11-05', 54000),

(104, 'Fatima Begum', 'fatima@example.com',
'Female', '1990-06-30', 62000);
```

Step 3: Use UNION to Combine Data

Let's combine the active and admin user names.

```
SELECT name FROM users

UNION

SELECT name FROM admin_users;
```

This returns a **single list of unique names** from both tables.

Advanced UNION Usage

CHAPTER 15

Multiple Columns and Sorting

UNION ALL Example

If you want to keep duplicate names (if any), use `UNION ALL`.

```
SELECT name FROM users

UNION ALL

SELECT name FROM admin_users;
```

Using More Than One Column

You can also select multiple columns as long as both `SELECT` queries return the **same number of columns with compatible types**.

```
SELECT name, salary FROM users

UNION

SELECT name, salary FROM admin_users;
```

Adding Separate Roles

```
SELECT name,
'User' AS role
FROM users

UNION

SELECT name,
'Admin' AS role
FROM admin_users;
```

UNION Summary

CHAPTER 15

Rules and Best Practices

Using Order By with UNION

```
SELECT name FROM users

UNION

SELECT name FROM admin_users

ORDER BY name;
```

Rules of UNION

1. The number of columns and their data types must match in all `SELECT` statements.
2. `UNION` removes duplicates by default.
3. `UNION ALL` keeps duplicates.

When to Use UNION

- When you have **two similar tables** (like current and archived data).
- When you need to **combine filtered results** from multiple sources.
- When performing **cross-category reporting**.

Summary

Operator	Behavior
<code>UNION</code>	Combines results, removes duplicates
<code>UNION ALL</code>	Combines results, keeps duplicates

Self JOIN in MySQL

CHAPTER 16

Joining a Table with Itself

A Self JOIN is a regular join, but the table is joined **with itself**.

This is useful when rows in the same table are related to each other. For example, when users refer other users, and we store the ID of the person who referred them in the same `users` table.

Step 1: Add a referred_by_id Column

We'll extend the existing `users` table to include a column called `referred_by_id`, which holds the `id` of the user who referred them.

```
ALTER TABLE users  
  
ADD COLUMN referred_by_id INT;
```

This column:

- Will be NULL for users who were not referred.
- Will contain the `id` of another user who referred them.

Step 2: Insert Referral Data (Optional)

Assuming `id = 1` is the first user, and referred others:

```
UPDATE users  
SET referred_by_id = 1  
WHERE id IN (2, 3);  
  
-- User 1 referred Users 2 and 3  
  
UPDATE users  
SET referred_by_id = 2  
WHERE id = 4;  
  
-- User 2 referred User 4
```


Using Self JOIN

CHAPTER 16

Getting Referrer Information

Step 3: Use a Self JOIN to Get Referrer Names

We want to get each user's name along with the name of the person who referred them.

```
SELECT

    a.id,

    a.name AS user_name,

    b.name AS referred_by

FROM users a

LEFT JOIN users b
ON a.referred_by_id = b.id;
```

Explanation

- `a` refers to the user being queried.
- `b` refers to the user who referred them.
- `LEFT JOIN` is used so that users with `NULL` in `referred_by_id` are also included.

Sample Output

id	user_name	referred_by
1	Aarav	NULL
2	Sneha	Aarav
3	Raj	Aarav
4	Fatima	Sneha

Self JOIN Summary

CHAPTER 16

Important Concepts and Best Practices

Summary

- Use Self JOIN when you need to **join a table with itself**.
- In referral-based relationships, store the referrer's `id` in the same table.
- Use aliases like `a` and `b` to differentiate the two instances of the same table.
- Self JOIN is commonly used in employee-manager structures, referral systems, and hierarchical data.



MySQL Views

CHAPTER 17

Working with Virtual Tables

A view in MySQL is a **virtual table** based on the result of a `SELECT` query. It does not store data itself. It always reflects the current data in the base tables.

Views are useful when:

- You want to simplify complex queries
- You want to reuse logic
- You want to hide certain columns from users
- You want a “live snapshot” of filtered data

Creating a View

Suppose we want a view that lists all users earning more than ₹70,000.

```
CREATE VIEW high_salary_users AS  
  
SELECT id, name, salary  
  
FROM users  
  
WHERE salary > 70000;
```

Querying the View

```
SELECT * FROM high_salary_users;
```

This will return all users from the `users` table where salary is above ₹70,000.

How Views Stay Updated

CHAPTER 17

Views Automatically Reflect Changes

Let's see what happens when the underlying data changes.

Step 1: View before update

```
SELECT * FROM high_salary_users;
```

Output:

id	name	salary
2	Sneha	75000
5	Fatima	80000

Step 2: Update a user's salary

```
UPDATE users  
  
SET salary = 72000  
  
WHERE name = 'Raj';
```

How Views Stay Updated

CHAPTER 17

Views Automatically Reflect Changes

Step 3: Query the view again

```
SELECT * FROM high_salary_users;
```

New Output:

id	name	salary
2	Sneha	75000
5	Fatima	80000
3	Raj	72000

Notice how Raj is now included in the view. The view did not need to be updated. That's because views always reflect live data from the original table.



Dropping a View

CHAPTER 17

Removing Existing Views

To remove a view:

```
DROP VIEW high_salary_users;
```

Summary

- Views act like saved SELECT queries
- Views are not duplicated data
- Changes to base tables are reflected automatically
- Great for simplifying complex queries or creating filtered access

MySQL Indexes

CHAPTER 18

Improving Query Performance

Indexes in MySQL are used to **speed up data retrieval**. They work like the index of a book, helping the database engine find rows faster, especially for **searches, filters, and joins**.

Viewing Indexes on a Table

To see the indexes on a table, use:

```
SHOW INDEXES FROM users;
```

This shows all the indexes currently defined on the `users` table, including the automatically created **primary key** index.

Creating a Single-Column Index

Suppose you're frequently searching users by their `email`. You can speed this up by indexing the `email` column.

```
CREATE INDEX idx_email ON users(email);
```

What this does:

- Creates an index named `idx_email`
- Improves performance of queries like:

```
SELECT * FROM users WHERE email = 'example@example.com';
```

Working with Multi-Column Indexes

CHAPTER 18

Optimizing Complex Searches

Important Notes

- Indexes **consume extra disk space**
- Indexes slow down `INSERT` , `UPDATE` , and `DELETE` operations slightly
- Use indexes only when needed (for columns used in `WHERE` , `JOIN` , and `ORDER BY`)

Creating a Multi-Column Index

If you often query users using both `gender` and `salary` , a multi-column index is more efficient than separate indexes.

```
CREATE INDEX idx_gender_salary ON users(gender, salary);
```

Usage Example:

```
SELECT * FROM users  
WHERE gender = 'Female' AND salary > 70000;
```

This query can take advantage of the combined index on `gender` and `salary` .

Index Order Matters

For a multi-column index on `(gender, salary)` :

- This works efficiently:

```
WHERE gender = 'Female' AND salary > 70000
```


Managing MySQL Indexes

CHAPTER 18

Best Practices and Removal

- But this may not use the index effectively:

```
WHERE salary > 70000
```

Because the first column in the index (`gender`) is missing in the filter.

Dropping an Index

To delete an index:

```
DROP INDEX idx_email ON users;
```

Summary

Feature	Description
SHOW INDEXES	View current indexes on a table
CREATE INDEX	Create single or multi-column indexes
DROP INDEX	Remove an index
Use when	Query performance on large tables is a concern
Avoid on	Columns that are rarely queried or always unique

Subqueries in MySQL

CHAPTER 19

Nested Queries and Dynamic Filtering

A subquery is a query nested inside another query. Subqueries are useful for breaking down complex problems into smaller parts.

They can be used in:

- `SELECT` statements
- `WHERE` clauses
- `FROM` clauses

Example Scenario: Salary Comparison

Suppose we want to find all users who **earn more than the average salary** of all users.

Scalar Subquery Example

This subquery returns a **single value**, the average salary. We then compare each user's salary with it.

```
SELECT id, name, salary
FROM users
WHERE salary > (
    SELECT AVG(salary) FROM users
);
```

Explanation:

- The inner query: `SELECT AVG(salary) FROM users` returns the average salary.
- The outer query selects all users with a salary **greater than that average**.

Using Subqueries

Filtering and Derived Values

Subquery with IN

Now let's say we want to find users who have been referred by someone who earns more than ₹75,000.

```
SELECT id, name, referred_by_id
FROM users
WHERE referred_by_id IN (
    SELECT id FROM users WHERE salary > 75000
);
```

Explanation:

- The inner query returns a list of user IDs who earn more than ₹75,000.
- The outer query selects users whose `referred_by_id` is in that list.

Other Places Subqueries Are Used

You can also use subqueries:

- Inside `SELECT` columns (called scalar subqueries)
- In the `FROM` clause to create derived tables

Example in `SELECT` :

```
SELECT name, salary,
    (SELECT AVG(salary) FROM users) AS average_salary
FROM users;
```

This shows each user's salary along with the overall average.

Subquery Summary

CHAPTER 19

Quick Recap of Common Use Cases

Summary

Subquery Type	Use Case
Scalar Subquery	Returns one value (e.g. AVG, MAX)
Subquery with IN	Returns multiple values
Subquery in SELECT	Shows related calculated value
Subquery in FROM	Acts as a virtual table

Subqueries are powerful tools when filtering based on computed or dynamic conditions.

GROUP BY and HAVING in MySQL

CHAPTER 20

Grouping and Filtering Aggregated Data

The `GROUP BY` clause is used to **group rows that have the same values** in specified columns. It is typically used with **aggregate functions** like `COUNT`, `SUM`, `AVG`, `MIN`, or `MAX`.

The `HAVING` clause is used to **filter groups** after aggregation. It is similar to how `WHERE` filters individual rows.

Example Table: `users`

Assume this is your `users` table:

id	name	gender	salary	referred_by_id
1	Aarav	Male	80000	NULL
2	Sneha	Female	75000	1
3	Raj	Male	72000	1
4	Fatima	Female	85000	2
5	Priya	Female	70000	NULL

GROUP BY Example: Average Salary by Gender

```
SELECT gender, AVG(salary) AS average_salary
FROM users
GROUP BY gender;
```

Explanation:

- This groups users by gender.
- Then calculates the average salary for each group.

Using GROUP BY and HAVING

CHAPTER 20

Aggregation and Group Filtering

GROUP BY with COUNT

Find how many users were referred by each user:

```
SELECT referred_by_id, COUNT(*) AS total_referred
FROM users
WHERE referred_by_id IS NOT NULL
GROUP BY referred_by_id;
```

Output:

referred_by_id	total_referred
1	2
2	1

HAVING Clause: Filtering Groups

Let's say we only want to show genders where the **average salary is greater than ₹75,000**.

```
SELECT gender, AVG(salary) AS avg_salary
FROM users
GROUP BY gender
HAVING AVG(salary) > 75000;
```

Why not WHERE ?

- WHERE is used before grouping.
- HAVING is used after groups are formed. It is the only way to filter aggregated values.

ROLLUP and Group Summaries

Advanced Aggregation Features

Another Example: Groups with More Than 1 Referral

```
SELECT referred_by_id, COUNT(*) AS total_referred
FROM users
WHERE referred_by_id IS NOT NULL
GROUP BY referred_by_id
HAVING COUNT(*) > 1;
```

Summary

Clause	Purpose	Can use aggregates?
WHERE	Filters rows before grouping	No
GROUP BY	Groups rows based on column values	N/A
HAVING	Filters groups after aggregation	Yes

Use `GROUP BY` to organize data, and `HAVING` to filter those groups based on aggregate conditions.

ROLLUP

To get subtotals and grand totals, you can use `ROLLUP` :

```
SELECT gender, COUNT(*) AS total_users
FROM users
GROUP BY gender WITH ROLLUP;
```

This gives you a count of users by gender, along with a grand total for all users.

Stored Procedures in MySQL

CHAPTER 21

Reusable SQL Blocks and Procedures

A stored procedure is a saved SQL block that can be executed later. It is useful for encapsulating logic that can be reused multiple times, such as queries, updates, or conditional operations.

Why Change the Delimiter?

By default, MySQL uses `;` to end SQL statements.

But when defining stored procedures, we use `;` inside the procedure as well. This can confuse MySQL. To avoid this, we **temporarily change the delimiter** (e.g. to `$$` or `//`) while creating the procedure.

Syntax for Creating a Stored Procedure

```
DELIMITER $$

CREATE PROCEDURE procedure_name()

BEGIN

    -- SQL statements go here

END$$

DELIMITER ;
```

After the procedure is created, we reset the delimiter back to `;`.

Creating Stored Procedures

CHAPTER 21

Using Input Parameters

Creating a Procedure with Input Parameters

Let's say you want to create a stored procedure that inserts a new user into the `users` table.

Example:

```
DELIMITER $$

CREATE PROCEDURE AddUser(

    IN p_name VARCHAR(100),

    IN p_email VARCHAR(100),

    IN p_gender ENUM('Male', 'Female', 'Other'),

    IN p_dob DATE,

    IN p_salary INT

)

BEGIN

    INSERT INTO users (name, email, gender, date_of_birth, salary)

    VALUES (p_name, p_email, p_gender, p_dob, p_salary);

END$$

DELIMITER ;
```

This creates a procedure named `AddUser` that accepts five input parameters.

Calling the Procedure

You can call the procedure using:

```
CALL AddUser('Kiran Sharma', 'kiran@example.com',
'Female', '1994-06-15', 72000);
```

This inserts the new user into the `users` table.

Managing Stored Procedures

CHAPTER 21

Viewing and Removing Procedures

Notes

- Input parameters are declared using the `IN` keyword.
- Stored procedures are stored in the database and can be reused.

Viewing Stored Procedures

```
SHOW PROCEDURE STATUS WHERE Db = 'startersql';
```

Dropping a Stored Procedure

```
DROP PROCEDURE IF EXISTS AddUser;
```

Summary

Command	Purpose
<code>DELIMITER \$\$</code>	Temporarily change statement delimiter
<code>CREATE PROCEDURE</code>	Defines a new stored procedure
<code>CALL procedure_name(...)</code>	Executes a stored procedure
<code>DROP PROCEDURE</code>	Removes an existing procedure

Triggers in MySQL

Automatically Executed Database Events

A trigger is a special type of stored program that is **automatically executed** when a specific event occurs in a table, such as `INSERT` , `UPDATE` , or `DELETE` .

Triggers are commonly used for:

- Logging changes
- Enforcing additional business rules
- Automatically updating related data

Basic Trigger Structure

```
CREATE TRIGGER trigger_name  
  
AFTER INSERT ON table_name  
  
FOR EACH ROW  
  
BEGIN  
  
    -- statements to execute  
  
END;
```

Triggers can be fired:

- `BEFORE` or `AFTER` an event
- On `INSERT` , `UPDATE` , or `DELETE`

Scenario: Log Every New User Insertion

Suppose we want to log every time a new user is inserted into the `users` table.

We'll create a separate table called `user_log` to store log entries.

Creating a Trigger

Logging Inserted Users

Step 1: Create the Log Table

```
CREATE TABLE user_log (  
  
    id INT AUTO_INCREMENT PRIMARY KEY,  
  
    user_id INT,  
  
    name VARCHAR(100),  
  
    created_on TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
  
);
```

Step 2: Create the Trigger

We now define a trigger that runs **after a new user is added**.

```
DELIMITER $$  
  
CREATE TRIGGER after_user_insert  
  
AFTER INSERT ON users  
  
FOR EACH ROW  
  
BEGIN  
  
    INSERT INTO user_log (user_id, name)  
  
    VALUES (NEW.id, NEW.name);  
  
END$$  
  
DELIMITER ;
```

Explanation:

- `AFTER INSERT` means the trigger fires after the user is inserted.
- `NEW` refers to the new row being added to the `users` table.
- We insert the new user's ID and name into the `user_log` table.

Managing Triggers

Testing and Removing Triggers

Step 3: Test the Trigger

```
CALL AddUser('Ritika Jain', 'ritika@example.com',  
'Female', '1996-03-12', 74000);
```

Now check the `user_log` table:

```
SELECT * FROM user_log;
```

You should see Ritika's info automatically logged.

Dropping a Trigger

If you need to remove a trigger:

```
DROP TRIGGER IF EXISTS after_user_insert;
```

Summary

Trigger Component	Description
<code>BEFORE / AFTER</code>	When the trigger runs
<code>INSERT / UPDATE / DELETE</code>	What kind of action triggers it
<code>NEW.column</code>	Refers to the new row (for <code>INSERT</code> , <code>UPDATE</code>)
<code>OLD.column</code>	Refers to the old row (for <code>UPDATE</code> , <code>DELETE</code>)
<code>FOR EACH ROW</code>	Executes for each affected row

More on MySQL

Additional Useful MySQL Features

This section covers some essential MySQL features and operators that help you write more powerful and flexible queries.

1. Logical Operators

Logical operators are used to combine multiple conditions in a `WHERE` clause.

Operator	Description	Example
AND	All conditions must be true	<code>salary > 50000 AND gender = 'Male'</code>
OR	At least one condition is true	<code>gender = 'Male' OR gender = 'Other'</code>
NOT	Reverses a condition	<code>NOT gender = 'Female'</code>

2. Add a Column to an Existing Table

Use `ALTER TABLE` to add a column:

```
ALTER TABLE users  
ADD COLUMN city VARCHAR(100);
```

This adds a new column named `city` to the `users` table.

Wildcards and LIMIT

Pattern Matching and Pagination

3. Wildcard Operators

Wildcards are used with the `LIKE` operator for pattern matching in text.

Wildcard	Description	Example
<code>%</code>	Matches any sequence	<code>WHERE name LIKE 'A%'</code> (starts with A)
<code>_</code>	Matches a single character	<code>WHERE name LIKE '_a%'</code> (second letter is 'a')

4. LIMIT with OFFSET

`LIMIT` is used to limit the number of rows returned. `OFFSET` skips rows before returning results.

```
SELECT * FROM users

ORDER BY id

LIMIT 5 OFFSET 10;
```

This skips the first 10 rows and returns the next 5.

Alternative syntax:

```
SELECT * FROM users

ORDER BY id

LIMIT 10, 5;
```

This also skips 10 and returns 5 (`LIMIT offset, count`).



DISTINCT and TRUNCATE

Filtering Unique Data and Clearing Tables

5. DISTINCT Keyword

`DISTINCT` is used to return **only unique values**.

```
SELECT DISTINCT gender FROM users;
```

Returns a list of unique gender values from the `users` table.

6. TRUNCATE Keyword

`TRUNCATE` removes all rows from a table, but **keeps the table structure**.

```
TRUNCATE TABLE users;
```

- Faster than `DELETE FROM users`
- Cannot be rolled back in most cases

7. CHANGE vs MODIFY Column

Both `CHANGE` and `MODIFY` are used to alter existing columns in a table.

CHANGE: Rename and change datatype

```
ALTER TABLE users  
  
CHANGE COLUMN city location VARCHAR(150);
```

This renames `city` to `location` and changes its type.

MODIFY: Only change datatype

```
ALTER TABLE users  
  
MODIFY COLUMN salary BIGINT;
```

This changes only the datatype of `salary`.